

Package ‘vbayesGP’

September 4, 2026

Type Package

Title Gaussian Variational Approximation to Gaussian Process Regression

Version 1.1.5

Date 2026-09-04

Author Seongil Jo [aut, cre],
Woojoo Lee [aut]

Maintainer Seongil Jo <bstatsjo@gmail.com>

Description Implements Gaussian variational approximation to Bayesian semiparametric regression with Gaussian process prior based on the Radial basis function (RBF) kernel. Consider the normal prior, the independent normal priors, or a continuous shrinkage prior on the lengthscale parameters of the RBF kernel.

License GPL-2

Imports Rcpp (>= 1.0.8), fields, ggplot2, MASS, ks, dplyr, magrittr,
tibble, tidyr

LinkingTo Rcpp, RcppArmadillo

Suggests bkmr

Depends R (>= 3.5.0)

RoxygenNote 7.3.3

Encoding UTF-8

LazyData true

LazyLoad yes

NeedsCompilation yes

Contents

computeOverallRisk	2
computePPIPs	3
computeSingVarInt	4
computeSingVarRisk	5
extractELBO	6
extractPostSamps	7
fitted.gpmim	7
fitted.gpr	8

gvaggpr	9
gvagpmim	13
gvagpr	16
metals	20
negbinomial	21
plot.gpmim	22
plot.gpr	22
plot.gprfitBivar	23
plot.gprfitBivarLevels	23
plot.gprfitUnivar	24
plot.risks	24
predict.gpmim	25
predict.gpr	26
predictorResponseBivar	27
predictorResponseBivarLevels	28
predictorResponseBivarPair	29
predictorResponseUnivar	30
print.gpmim	32
print.gpr	32
summary.gpmim	33
summary.gpr	33

Index 35

computeOverallRisk	<i>Calculate Overall Risk Summaries</i>
--------------------	---

Description

Compare estimated f function when all predictors are at a particular quantile to when all are at a second fixed quantile

Usage

```
computeOverallRisk(
  object,
  qs = seq(0.25, 0.75, by = 0.05),
  q.fixed = 0.5,
  nsamples = 1000,
  ...
)
```

Arguments

object	an object of class gpr obtained from gvagpr function
qs	vector of quantiles at which to calculate the overall risk summary
q.fixed	a second quantile at which to compare the estimated f function
nsamples	(positive integer), number of posterior samples to draw and save, defaults to 1000
...	unused argument

Value

a data frame containing the (posterior mean) estimate and posterior standard deviation of the overall risk measures

References

Bobb J (2023). `_bkmr: Bayesian Kernel Machine Regression_`. R package version 0.2.2, <<https://github.com/jenfb/bkmr>>

See Also

`gvagpr`

Examples

```
## Not run:
## First generate dataset
set.seed(111)
dat <- bkmr::SimData(n = 50, M = 4)
y <- dat$y
Z <- dat$Z
X <- dat$X

set.seed(111)
priors <- list(lengthscale = 'horseshoe')
fout <- vbayesGP::gvagpr(y = y, Z = Z, X = X, priors = priors, covstr = 'diagonal')
risks.overall <- vbayesGP::computeOverallRisk(fout, qs = seq(0.25, 0.75, by = 0.05), q.fixed = 0.5)
plot(risk.overall)
## End(Not run)
```

computePPIPs

Compute pseudo posterior inclusion probabilities (PPIPs) from VGPR model or VGPMIM fits

Description

Compute pseudo posterior inclusion probabilities (PPIPs) using Sequential-2-Means from Variational Gaussian Process Regression (VGPR) model fit or Variational Gaussian Process Multiple Index Model (VGPMIM) fit

Usage

```
computePPIPs(object, nsamples = 1000, ntuning = 10)
```

Arguments

<code>object</code>	an object of class <code>gpr</code> or <code>gpmim</code>
<code>nsamples</code>	(positive integer), number of posterior samples to draw and save, defaults to 1000
<code>ntuning</code>	the number of chosen values of the tuning parameter for variable selection

Value

a data frame including the variable-specific PPIPs for VGPR fit with horseshoe prior

Author(s)

Seongil Jo

References

Li, H. and Pati, D. (2017). "Variable Selection Using Shrinkage Priors", Computational Statistics and Data Analysis, 107, 107-119.

See Also

gvagpr, gvaggpr, gvagpmim

computeSingVarInt

*Single Variable Interaction Summaries***Description**

Compare the single-predictor health risks when all of the other predictors in Z are fixed to their a specific quantile to when all of the other predictors in Z are fixed to their a second specific quantile.

Usage

```
computeSingVarInt(
  object,
  which.z,
  qs.diff = c(0.25, 0.75),
  qs.fixed = c(0.25, 0.75),
  nsamples = 1000,
  ...
)
```

Arguments

object	an object of class gpr obtained from gvagpr function
which.z	vector indicating which variables (columns of Z) for which the summary should be computed
qs.diff	vector indicating the two quantiles at which to compute the single-predictor risk summary
qs.fixed	vector indicating the two quantiles at which to fix all of the remaining exposures in Z
nsamples	(positive integer), number of posterior samples to draw and save, defaults to 1000
...	unused argument

Value

a data frame containing the (posterior mean) estimate and posterior standard deviation of the single-predictor risk measures

Examples

```
## Not run:
## First generate dataset
set.seed(111)
dat <- bkmr::SimData(n = 50, M = 4)
y <- dat$y
Z <- dat$Z
X <- dat$X

set.seed(111)
priors <- list(lengthscale = 'horseshoe')
fout <- vbayesGP::gvagpr(y = y, Z = Z, X = X, priors = priors, covstr = 'diagonal')
risks.int <- computeSingVarInt(fout)
plot(risks.int)
## End(Not run)
```

computeSingVarRisk	<i>Single Variable Risk Summaries</i>
--------------------	---------------------------------------

Description

Compute summaries of the risks associated with a change in a single variable in Z from a single level (quantile) to a second level (quantile), for the other variables in Z fixed to a specific level (quantile)

Usage

```
computeSingVarRisk(
  object,
  which.z,
  qs.diff = c(0.25, 0.75),
  q.fixed = c(0.25, 0.5, 0.75),
  nsamples = 1000,
  ...
)
```

Arguments

object	an object of class gpr obtained from gvagpr function.
which.z	vector indicating which variables (columns of Z) for which the summary should be computed
qs.diff	vector indicating the two quantiles q_{-1} and q_{-2} at which to compute $f(z_{-}\{q_2\}) - f(z_{-}\{q_1\})$
q.fixed	vector of quantiles at which to fix the remaining predictors in Z
nsamples	(positive integer), number of posterior samples to draw and save, defaults to 1000
...	unused argument

Value

a data frame containing the (posterior mean) estimate and posterior standard deviation of the single-predictor risk measures

References

Bobb J (2023). `_bkmr: Bayesian Kernel Machine Regression_`. R package version 0.2.2, <<https://github.com/jenfb/bkmr>>

Examples

```
## Not run:
## First generate dataset
set.seed(111)
dat <- bkmr::SimData(n = 50, M = 4)
y <- dat$y
Z <- dat$Z
X <- dat$X

set.seed(111)
fout <- vbayesGP::gvagpr(y = y, Z = Z, X = X, priors = priors, covstr = 'diagonal')
risks.singvar <- computeSingVarRisk(fout)
plot(risk.singvar)
## End(Not run)
```

extractELBO

Extract ELBO from gpr, ggpr, or gpmim model fits

Description

Compute the expected lower bound (ELBO) using the posterior samples for class "gpr" of "gpmim"

Usage

```
extractELBO(object, nsamples = 1000)
```

Arguments

object	an object of class gpr or gpmim.
nsamples	(positive integer), number of posterior samples to draw and save, defaults to 1000.

Author(s)

Seongil Jo

See Also

gvagpr, gvaggpr, gvagpmmim

extractPostSamps	<i>Extract Posterior Samples from gpr, ggpr, or gpmim model fits</i>
------------------	--

Description

Generate the posterior samples for class "gpr" or "gpmim"

Usage

```
extractPostSamps(object, nsamples = 1000)
```

Arguments

object	an object of class gpr or gpmim.
nsamples	(positive integer), number of posterior samples to draw and save, defaults to 1000.

Value

a data frame including posterior samples for all parameters. If family is gaussian, the posterior samples contain β , σ^2 , λ_f , and γ . If family is not gaussian, the posterior samples also contain $f_i, i = 1, \dots, N$. If `object$id` is not NULL, the data frame also includes $b_i, i = 1, \dots, N$.

Author(s)

Seongil Jo

See Also

gvagpr, gvaggpr, gvagpmmim

fitted.gpmim	<i>Extract GPMIM Fitted Values</i>
--------------	------------------------------------

Description

fitted is a generic function which extracts fitted values of nonparametric part from an object of class "gpmim"

Usage

```
## S3 method for class 'gpmim'
fitted(object, nsamples = 1000, cov = TRUE, ...)
```

Arguments

object	an object of class gpmim.
nsamples	(positive integer) number of posterior samples. Default value is 1000.
cov	logical. If TRUE (the default) the posterior covariance matrix of the nonparametric part is returned as well. This matrix is n by n, so for large samples it may dominate the memory used by the call; set cov = FALSE to return the fitted means only.
...	unused argument.

Value

fmean	posterior mean of nonparametric part.
fcov	posterior variance of nonparametric part.
an object of class "gprfit", which has the associated method:	
* plot (i.e., plot.gprfit)	

Author(s)

Seongil Jo

See Also

gvagpmim

fitted.gpr

Extract GPR and GGPR Model Fitted Values

Description

fitted is a generic function which extracts fitted values of nonparametric part from an object of class "gpr"

Usage

```
## S3 method for class 'gpr'
fitted(object, nsamples = 1000, cov = TRUE, ...)
```

Arguments

object	an object of class gpr.
nsamples	(positive integer) number of posterior samples. Default value is 1000.
cov	logical. If TRUE (the default) the posterior covariance matrix of the nonparametric part is returned as well. This matrix is n by n, so for large samples it may dominate the memory used by the call; set cov = FALSE to return the fitted means only.
...	unused argument.

Value

fmean posterior mean of nonparametric part.
 fcov posterior variance of nonparametric part, or NULL if cov = FALSE.
 an object of class "gprfit", which has the associated method:
 * plot (i.e., plot.gprfit)

Author(s)

Seongil Jo

See Also

gvagpr, gvaggpr

gvaggpr

Gaussian Variational Approximation to Generalized Gaussian Process Regression

Description

Fits the Bayesian kernel machine regression for generalized linear model using Gaussian variational approximation algorithm.

Usage

```
gvaggpr(
  y,
  X,
  Z,
  id = NULL,
  random.slope = NULL,
  family = binomial(),
  priors = list(),
  inits = list(),
  covstr = c("diagonal", "fullrank", "blockdiag", "sparseprec"),
  control = list(),
  minibatch = TRUE,
  verbose = TRUE,
  seed = sample.int(.Machine$integer.max, 1)
)
```

Arguments

y a vector of response of length n.
 X an n-by-p matrix of covariates for parametric term. Should not contain an intercept.
 Z an n-by-M matrix of predictor variables to be included in nonparametric part.
 id optional vector (of length n) of grouping factors for fitting a model with random effects (including both a random intercept and a random slope). If NULL then no random effects will be included.

<code>random.slope</code>	a column index of the matrix (X) including covariates for random slope. If NULL and <code>id</code> is given, the model considers the random intercept only.
<code>family</code>	a description of the error distribution and link function. Either <code>binomial()</code> with the logit link, or <code>MASS::negative.binomial(theta =)</code> with the log link, in which case the linear predictor is the log mean, as in <code>MASS::glm.nb</code> . The dispersion is always estimated; the value given to <code>theta</code> is its starting value.
<code>priors</code>	a list giving the prior information. The list includes the following parameters (with default values in parentheses): <code>alam</code> (1) and <code>blam</code> (0.1) giving the hyper parameters for λ_f , <code>tau0</code> (1) giving the hyper parameters of the beta-prime prior, <code>s_theta</code> (1) giving the scale of the half-normal prior placed on $1/\sqrt{r}$, the standard deviation of the multiplicative random effect in the Poisson–gamma representation of the negative binomial. The element <code>lengthscale</code> selects the prior on the inverse length-scale parameters and is one of 'normal' (a single length-scale shared by all exposures), 'indep.normal' (one length-scale per exposure, the default when <code>ncol(Z) > 1</code>) or 'betaprime' (one per exposure with a continuous shrinkage prior, which performs variable selection). 'shrinkage' is accepted as a deprecated alias for 'betaprime'.
<code>inits</code>	a named list of starting values for the variational means, given by name rather than by position, so that the order in which the blocks are stored does not matter. Every element is given on the natural scale and is logged internally where the parameter is positive; a scalar is recycled over a block. The names are <code>beta</code> (the regression coefficients), <code>lengthscale</code> (the inverse length-scales γ), <code>lambda</code> (the scale of the Gaussian process), <code>taur</code> (the global shrinkage scale, for <code>lengthscale = "betaprime"</code>), and, for the negative binomial family, <code>theta</code> (the dispersion). Anything not named starts at zero on the working scale. The inducing values and the random-effect covariance are set internally and are not exposed.
<code>covstr</code>	Either "diagonal" (the default), "fullrank", "blockdiag", or "sparseprec", indicating which covariance structure of variational distribution is used. The "diagonal" option uses a fully factorized Gaussian for the approximation, and is available for every model. For a model without random effects, the "fullrank" option uses a Gaussian with an unrestricted lower triangular Cholesky factor. For the mixed effect models the parameter vector is partitioned by cluster and an unrestricted factor is not offered: the "blockdiag" option uses a Gaussian whose covariance matrix is block diagonal, with one block per cluster and one block for the parameters shared by all clusters, and the "sparseprec" option uses a Gaussian with a sparse precision matrix, which adds the dependence between each cluster and the shared parameters. An option that does not apply to the model at hand falls back to the closest one that does, with a warning: "fullrank" becomes "blockdiag" when random effects are present, and "blockdiag" and "sparseprec" become "fullrank" when they are not.
<code>control</code>	a named list of parameters to control the algorithm's behavior. The list includes the following parameters (with default values in parentheses): <code>max_iter</code> (100000) giving the maximum number of iterations, <code>nbatch</code> (<code>n / 100</code>) giving the number of mini-batches when <code>minibatch = TRUE</code> , <code>alpha0_mu</code> (1e-3) and <code>alpha0_L</code> (1e-4) giving the base step sizes for the mean and Cholesky parameters, respectively, <code>beta</code> (0.9) giving the exponential moving-average coefficient used to smooth the (normalized) natural-gradient direction, <code>nws</code> (2500) giving rolling window size for calculating the moving average of the lower bounds, <code>nsp</code> (100) giving the maximum patience parameter, <code>ninpts</code> (50) giving the number of inducing points for sparse Gaussian process prior. <code>knots</code> (NULL) giving the inducing points for sparse Gaussian process prior.

minibatch	TRUE or FALSE: If TRUE, nbatch (the number of batch) should be given in control argument and max_iter denotes the number of epoch. Default value is $n/100$. Note. this option is not applicable for the "sparseprec" option of covstr and the random effects models.
verbose	TRUE or FALSE: flag indicating whether to print intermediate diagnostic information during the model fitting.
seed	The seed for random number generation. The default is generated from 1 to the maximum integer supported by R on the machine.

Details

Jo et al. (2026+) proposed the Bayesian semiparametric generalized linear model with sparse Gaussian process prior based on the Radial basis function (RBF) kernel:

$$\begin{aligned}
 y_i &\sim \text{Bern}(p_i), \log(p_i/(1+p_i)) = x_i^\top \beta + f(z_i), \\
 f &= (f(z_1), \dots, f(z_D))^\top \sim GP(K^{-1} \bar{f}, \sigma^2 \lambda_f \tilde{K}), \quad z_i = (z_{i1}, \dots, z_{iM})^\top, \\
 \bar{f} &= (f(\bar{z}_1), \dots, f(\bar{z}_{\bar{D}}))^\top \sim GP(0, K_{\bar{D}}),
 \end{aligned}$$

where $\bar{D} \ll D$, $K = K_D K_{DD}^{-1}$, $\tilde{K} = K_D - K_{D\bar{D}} K_{DD}^{-1} K_{\bar{D}D}$, and the kernel denotes the RBF kernel given as

1) Equal lengthscale parameter:

$$K_D = \left(\exp \left(-\gamma \sum_{m=1}^M \|z_i - z_j\|^2 \right) \right)_{i,j=1}^D$$

2) Varying lengthscale parameters:

$$K_D = \left(\exp \left(- \sum_{m=1}^M \gamma_m \|z_i - z_j\|^2 \right) \right)_{i,j=1}^D$$

For the parameters, the following priors are used:

$$\pi(\beta) \propto 1,$$

$$\pi(\lambda_f) = \text{Gamma}(a_\lambda, b_\lambda),$$

1) Normal prior:

$$\pi(\gamma) = N_+(0, \tau_0^2)$$

2) Independent Normal priors:

$$\pi(\gamma_m) = N_+(0, \tau_0^2), \quad m = 1, \dots, M$$

3) Beta-prime (continuous shrinkage) prior:

$$\gamma_m \mid \tau_\gamma \sim \text{BetaPrime} \left(\frac{1}{2}, \frac{1}{2}; \text{scale} = \tau_\gamma^2 \right), \quad m = 1, \dots, M,$$

$$\pi(\gamma_m \mid \tau_\gamma) \propto \gamma_m^{-1/2} (1 + \gamma_m / \tau_\gamma^2)^{-1}, \quad \gamma_m > 0,$$

$$\tau_\gamma \sim \mathcal{C}^+(0, \tau_0),$$

where τ_0 is a positive constant specified by the user.

For more details, see Jo and Lee (2026+).

Value

an object of class "gpr", which has the associated methods:

- * extractELBO
- * fitted (i.e., fitted.gpr)
- * summary (i.e., summary.gpr)
- * predict (i.e., predict.gpr)
- * plot (i.e., plot.gpr)

Author(s)

Seongil Jo and Woojoo Lee

References

- Jo, S., and Lee, W. (2025+), "Gaussian variational approximation to Bayesian kernel machine regression for Logistic models", *preprint*.
- Jo, S., Ye, S., Hanh, G., and Lee, W. (2026), "Accelerated Bayesian Kernel Machine Regression: A Gaussian Variational Approximation with the Horseshoe Prior", *Statistics and Computing*.
- Polson, N. G., Scott, J. G., and Windle, J. (2013), "Bayesian Inference for Logistic Models using Polya-Gamma Latent Variables", *Journal of the American Statistical Association*, 108, 1339-1349.
- Bobb, J. F., Valeri, L., Claus, H. B., Christiani, D. C., Wright, R. O., Mazumdar, M., Godleski, J. J., and Coull, B. A. (2015). "Bayesian Kernel Machine Regression for Estimating the Health Effects of Multi-Pollutant Mixtures", *Biostatistics*, 16, 493-508.
- Tan, L. S. L. and Nott D. J. (2018), "Gaussian variational approximation with sparse precision matrices", *Statistics and Computing*, 28, 259-275.
- Tan, L. S. L. (2025), "Analytic natural gradient updates for Cholesky factor in Gaussian variational approximation", *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 00, 1-27.

See Also

extractELBO, fitted.gpr, predict.gpr, plot.gpr, summary.gpr

Examples

```
## Not run:
sdat <- bkmr::SimData()
y <- sdat$y
X <- sdat$X
Z <- sdat$Z

fout <- vbayesGP::gvaggpr(y, X, Z, priors = list(lengthscale = 'normal'), family = binomial())
plot(fout)
summary(fout)
vbayesGP::extractELBO(fout) # ELBO
## End(Not run)
```

gvagpmim

*Gaussian Variational Approximation to Gaussian Process Multiple Index Model***Description**

Fits the Bayesian kernel machine multiple index model using Gaussian variational approximation algorithm.

Usage

```
gvagpmim(
  y,
  X,
  Z,
  m.index = NULL,
  id = NULL,
  random.slope = NULL,
  priors = list(),
  covstr = c("diagonal", "fullrank", "blockdiag", "sparseprec"),
  control = list(),
  minibatch = FALSE,
  verbose = TRUE,
  seed = sample.int(.Machine$integer.max, 1)
)
```

Arguments

y	a vector of response of length n.
X	a matrix of covariates for parametric term. Should not contain an intercept.
Z	a matrix or a list of predictor variables to be included in nonparametric part.
m.index	a list of column indices grouping the predictor variables in nonparametric part. If Z is a matrix, m.index must be given.
id	optional vector (of length n) of grouping factors for fitting a model with random effects (including both a random intercept and a random slope). If NULL then no random effects will be included.
random.slope	a column index of the matrix (X) including covariates for random slope. If NULL and id is given, the model considers the random intercept only.
priors	a list giving the prior information. The list includes the following parameters (with default values in parentheses): asig (0.001) and bsig (0.001) giving the hyper parameters for σ^2 , alam (0.1) and blam (0.01) giving the hyper parameters for λ_f , lam0 (1) and tau0 (1) giving the hyper parameters of the horseshoe prior. The element lengthscale selects the prior on the reparameterized index coefficients and is one of 'indep.normal' (independent normal priors, the default) or 'horseshoe' (a continuous shrinkage prior, which performs variable selection). There is no isotropic option here, since the index parameterization already carries one component per exposure. 'normal' is accepted as a deprecated alias for 'indep.normal'.

covstr	Either "diagonal" (the default), "fullrank", "blockdiag", or "sparseprec", indicating which covariance structure of variational distribution is used. The "diagonal" option uses a fully factorized Gaussian for the approximation, and is available for every model. For a model without random effects, the "fullrank" option uses a Gaussian with an unrestricted lower triangular Cholesky factor. For the mixed effect models the parameter vector is partitioned by cluster and an unrestricted factor is not offered: the "blockdiag" option uses a Gaussian whose covariance matrix is block diagonal, with one block per cluster and one block for the parameters shared by all clusters, and the "sparseprec" option uses a Gaussian with a sparse precision matrix, which adds the dependence between each cluster and the shared parameters. An option that does not apply to the model at hand falls back to the closest one that does, with a warning: "fullrank" becomes "blockdiag" when random effects are present, and "blockdiag" and "sparseprec" become "fullrank" when they are not.
control	a named list of parameters to control the algorithm's behavior. The list includes the following parameters (with default values in parentheses): max_iter (100000) giving the maximum number of iterations, rho (0.95) giving the decaying constant, eps (1e-6) giving the small positive constant added to ensure the denominator of the step size is positive and the initial step size is nonzero, nws (2500) giving rolling window size for calculating the moving average of the lower bounds, nsp (100) giving the maximum patience parameter.
minibatch	TRUE or FALSE: If TRUE, nbatch (the number of batch) should be given in control argument and max_iter denotes the number of epoch. Default value is $n/100$. Note. this option is not applicable for the "sparseprec" option of covstr and the random effects models.
verbose	TRUE or FALSE: flag indicating whether to print intermediate diagnostic information during the model fitting.
seed	The seed for random number generation. The default is generated from 1 to the maximum integer supported by R on the machine.

Details

Jo et al. (2026) proposed the Bayesian semiparametric regression model with Gaussian process prior based on the Radial basis function (RBF) kernel:

$$y_i = x_i^\top \beta + f^M(z_{i1}^\top \alpha_1, \dots, z_{iM}^\top \alpha_M) + \epsilon_i, \quad \epsilon_i \stackrel{iid}{\sim} N(0, \sigma^2),$$

$$f_i^M = f^M(z_{i1}^\top \alpha_1, \dots, z_{iM}^\top \alpha_M),$$

$$f^M = (f_1^M, \dots, f_D^M)^\top \sim GP(0, \sigma^2 \lambda_f K_D),$$

where K_D denotes the RBF kernel given as

Varying lengthscale parameters:

$$K_D = \left(\exp \left(- \sum_{m=1}^M ((z_{im} - z_{jm})^\top D_m^{-1} \gamma_m)^2 \right) \right)_{i,j=1}^D, \quad \gamma_m = (\gamma_{m1}, \dots, \gamma_{mL_m})^\top,$$

where D_m is a non-singular matrix for constraints.

Jo et al. (2026) also proposed the Bayesian semiparametric mixed effects regression model with Gaussian process prior based on the Radial basis function (RBF) kernel:

$$y_{ir} = x_{ir}^\top \beta + f^M(z_{ir1}^\top \alpha_1, \dots, z_{irM}^\top \alpha_M) + u_{ir}^\top b_i + \epsilon_{ir}, \quad \epsilon_{ir} \stackrel{iid}{\sim} N(0, \sigma^2),$$

$$b_i = (b_{i1}, \dots, b_{iq})^\top \stackrel{iid}{\sim} N(0, \Sigma_b).$$

For the parameters, the following priors are used:

$$\begin{aligned}\pi(\beta) &\propto 1, \\ \pi(\sigma^{-2}) &= \text{Gamma}(a_\sigma, b_\sigma), \\ \pi(\lambda_f) &= \text{Gamma}(a_\lambda, b_\lambda),\end{aligned}$$

1) Independent Normal priors:

$$\begin{aligned}\pi(\gamma_{ml}) &= N(0, \tau_0^2), \quad m = 1, \dots, M, \quad l = 1, \dots, L_m - 1 \\ \pi(\gamma_{mL_m}) &= N_+(0, \tau_0^2), \quad m = 1, \dots, M\end{aligned}$$

2) Horseshoe prior:

$$\begin{aligned}\pi(\gamma_{ml} \mid \lambda_{ml}, \tau_m, \tau_\gamma) &= N(0, \lambda_{ml}^2 \tau_m^2 \tau_\gamma^2), \quad m = 1, \dots, M, \quad l = 1, \dots, L_m - 1 \\ \pi(\gamma_{mL_m} \mid \lambda_{mL_m}, \tau_m, \tau_\gamma) &= N_+(0, \lambda_{mL_m}^2 \tau_m^2 \tau_\gamma^2), \quad m = 1, \dots, M \\ \pi(\lambda_{ml}) &= C_+(0, \lambda_0), \quad m = 1, \dots, M, \quad l = 1, \dots, L_m, \\ \pi(\tau_m) &= C_+(0, \tau_0), \quad m = 1, \dots, M, \\ \pi(\tau_\gamma) &= C_+(0, \tau_{\gamma,0}),\end{aligned}$$

where $a_\sigma, b_\sigma, a_\lambda, b_\lambda, \lambda_0, \tau_0$ and $\tau_{\gamma,0}$ are positive constants specified by users.

For more details, see Jo et al. (2026).

Value

an object of class "gpmim", which has the associated methods:

- * extractELBO
- * fitted (i.e., fitted.gpmim)
- * summary (i.e., summary.gpmim)
- * predict (i.e., predict.gpmim)
- * plot (i.e., plot.gpmim)

Author(s)

Seongil Jo and Woojoo Lee

References

- Jo, S., Ye, S., Hanh, G., and Lee, W. (2026), "Accelerated Bayesian Kernel Machine Regression: A Gaussian Variational Approximation with the Horseshoe Prior", *Statistics and Computing*, 36:79.
- Titsias, M. K. and Lazaro-Gredilla, M. (2014), "Doubly stochastic variational Bayes for non-conjugate inference", *Proceedings of the 31st ICML*.
- Bobb, J. F., Valeri, L., Claus, H. B., Christiani, D. C., Wright, R. O., Mazumdar, M., Godleski, J. J., and Coull, B. A. (2015). "Bayesian Kernel Machine Regression for Estimating the Health Effects of Multi-Pollutant Mixtures", *Biostatistics*, 16, 493-508.
- Chen, H., Zheng, L., Kontai, R. A., and Raskutti, G. (2022), "Gaussian process parameter estimation using mini-batch stochastic gradient descent: convergence guarantees and empirical benefits", *Journal of Machine Learning Research*, 23, 1-59.
- McGee, G., Wilson, A., Webster, T. F., Coull, B. A. (2021), "Bayesian multiple index models for environmental mixtures", *Biometrics*, 1-13.

See Also

extractELBO, fitted.gpmim, predict.gpmim, plot.gpmim, summary.gpmim

Examples

```
## Not run:
sdat <- bkmr::SimData(M = 13)
y <- sdat$y
X <- sdat$X
Z <- sdat$Z
m.index <- list(1:5, 6:13)

fout <- vbayesGP::gvagpmim(y, X, Z, m.index, priors = list(lengthscale = 'indep.normal'))
plot(fout)
summary(fout)
vbayesGP::extractELBO(fout) # ELBO
## End(Not run)
```

gvagpr

Gaussian Variational Approximation to Gaussian Process Regression

Description

Fits the Bayesian kernel machine regression using Gaussian variational approximation algorithm.

Usage

```
gvagpr(
  y,
  X,
  Z,
  id = NULL,
  random.slope = NULL,
  priors = list(),
  covstr = c("diagonal", "fullrank", "blockdiag", "sparseprec"),
  inits = list(),
  control = list(),
  minibatch = FALSE,
  verbose = TRUE,
  seed = sample.int(.Machine$integer.max, 1)
)
```

Arguments

y	a vector of response of length n.
X	an n-by-p matrix of covariates for parametric term. Should not contain an intercept.
Z	an n-by-M matrix of predictor variables to be included in nonparametric part.
id	optional vector (of length n) of grouping factors for fitting a model with random effects (including both a random intercept and a random slope). If NULL then no random effects will be included.

random.slope	a column index of the matrix (X) including covariates for random slope. If NULL and id is given, the model considers the random intercept only.
priors	a list giving the prior information. The list includes the following parameters (with default values in parentheses): asig (0.001) and bsig (0.001) giving the hyper parameters for σ^2 , alam (0.1) and blam (0.01) giving the hyper parameters for λ_f , lam0 (1) and tau0 (1) giving the hyper parameters of the horseshoe prior. The element lengthscale selects the prior on the inverse length-scale parameters and is one of 'normal' (a single length-scale shared by all exposures), 'indep.normal' (one length-scale per exposure, the default when ncol(Z) > 1) or 'horseshoe' (one per exposure with a continuous shrinkage prior, which performs variable selection).
covstr	Either "diagonal" (the default), "fullrank", "blockdiag", or "sparseprec", indicating which covariance structure of variational distribution is used. The "diagonal" option uses a fully factorized Gaussian for the approximation, and is available for every model. For a model without random effects, the "fullrank" option uses a Gaussian with an unrestricted lower triangular Cholesky factor. For the mixed effect models the parameter vector is partitioned by cluster and an unrestricted factor is not offered: the "blockdiag" option uses a Gaussian whose covariance matrix is block diagonal, with one block per cluster and one block for the parameters shared by all clusters, and the "sparseprec" option uses a Gaussian with a sparse precision matrix, which adds the dependence between each cluster and the shared parameters. An option that does not apply to the model at hand falls back to the closest one that does, with a warning: "fullrank" becomes "blockdiag" when random effects are present, and "blockdiag" and "sparseprec" become "fullrank" when they are not.
inits	a named list of starting values for the variational means, given by name rather than by position. Every element is on the natural scale and is logged internally where the parameter is positive; a scalar is recycled over a block. The names are beta (the regression coefficients), lengthscale (the inverse length-scales γ), lambda (the scale of the Gaussian process), sigsq.eps (the error variance), and, for lengthscale = "horseshoe", tau and taur (the local and global shrinkage scales). Anything not named starts at zero on the working scale. The random-effect covariance is set internally and is not exposed.
control	a named list of parameters to control the algorithm's behavior. The list includes the following parameters (with default values in parentheses): max_iter (100000) giving the maximum number of iterations, rho (0.95) giving the decaying constant, eps (1e-6) giving the small positive constant added to ensure the denominator of the step size is positive and the initial step size is nonzero, nws (2500) giving rolling window size for calculating the moving average of the lower bounds, nsp (100) giving the maximum patience parameter.
minibatch	TRUE or FALSE: If TRUE, nbatch (the number of batch) should be given in control argument and max_iter denotes the number of epoch. Default value is $n/100$. Note. this option is not applicable for the "sparseprec" option of covstr and the random effects models.
verbose	TRUE or FALSE: flag indicating whether to print intermediate diagnostic information during the model fitting.
seed	The seed for random number generation. The default is generated from 1 to the maximum integer supported by R on the machine.

Details

Jo et al. (2026) proposed the Bayesian semiparametric regression model with Gaussian process prior based on the Radial basis function (RBF) kernel:

$$y_i = x_i^\top \beta + f(z_i) + \epsilon_i, \quad \epsilon_i \stackrel{iid}{\sim} N(0, \sigma^2),$$

$$f = (f(z_1), \dots, f(z_D))^\top \sim GP(0, \sigma^2 \lambda_f K_D), \quad z_i = (z_{i1}, \dots, z_{iM})^\top,$$

where K_D denotes the RBF kernel given as

1) Equal lengthscale parameter:

$$K_D = \left(\exp \left(-\gamma \sum_{m=1}^M \|z_i - z_j\|^2 \right) \right)_{i,j=1}^D$$

2) Varying lengthscale parameters:

$$K_D = \left(\exp \left(-\sum_{m=1}^M \gamma_m \|z_{im} - z_{jm}\|^2 \right) \right)_{i,j=1}^D$$

Jo et al. (2026) also proposed the Bayesian semiparametric mixed effects regression model with Gaussian process prior based on the Radial basis function (RBF) kernel:

$$y_{ir} = x_{ir}^\top \beta + f(z_{ir}) + u_{ir}^\top b_i + \epsilon_{ir}, \quad \epsilon_{ir} \stackrel{iid}{\sim} N(0, \sigma^2),$$

$$b_i = (b_{i1}, \dots, b_{iq})^\top \stackrel{iid}{\sim} N(0, \Sigma_b).$$

For the parameters, the following priors are used:

$$\pi(\beta) \propto 1,$$

$$\pi(\sigma^{-2}) = \text{Gamma}(a_\sigma, b_\sigma),$$

$$\pi(\lambda_f) = \text{Gamma}(a_\lambda, b_\lambda),$$

1) Normal prior:

$$\pi(\gamma) = N_+(0, \tau_0^2)$$

2) Independent Normal priors:

$$\pi(\gamma_m) = N_+(0, \tau_0^2), \quad m = 1, \dots, M$$

3) Horseshoe prior:

$$\pi(\gamma_m \mid \lambda_m, \tau_\gamma) = N_+(0, \lambda_m^2 \tau_\gamma^2), \quad m = 1, \dots, M$$

$$\pi(\lambda_m) = C_+(0, \lambda_0), \quad m = 1, \dots, M$$

$$\pi(\tau_\gamma) = C_+(0, \tau_0),$$

where $a_\sigma, b_\sigma, a_\lambda, b_\lambda, \lambda_0$ and τ_0 are positive constants specified by users.

For more details, see Jo et al. (2026).

Value

an object of class "gpr", which has the associated methods:

- * extractELBO
- * fitted (i.e., fitted.gpr)
- * summary (i.e., summary.gpr)
- * predict (i.e., predict.gpr)
- * plot (i.e., plot.gpr)

Author(s)

Seongil Jo and Woojoo Lee

References

- Jo, S., Ye, S., Hanh, G., and Lee, W. (2026), "Accelerated Bayesian Kernel Machine Regression: A Gaussian Variational Approximation with the Horseshoe Prior", *Statistics and Computing*, 36:79.
- Titsias, M. K. and Lazaro-Gredilla, M. (2014), "Doubly stochastic variational Bayes for non-conjugate inference", *Proceedings of the 31st ICML*.
- Bobb, J. F., Valeri, L., Claus, H. B., Christiani, D. C., Wright, R. O., Mazumdar, M., Godleski, J. J., and Coull, B. A. (2015). "Bayesian Kernel Machine Regression for Estimating the Health Effects of Multi-Pollutant Mixtures", *Biostatistics*, 16, 493-508.
- Tan, L. S. L. and Nott D. J. (2018), "Gaussian variational approximation with sparse precision matrices", *Statistics and Computing*, 28, 259-275.
- Chen, H., Zheng, L., Kontai, R. A., and Raskutti, G. (2022), "Gaussian process parameter estimation using mini-batch stochastic gradient descent: convergence guarantees and empirical benefits", *Journal of Machine Learning Research*, 23, 1-59.
- Tan, L. S. L. and Nott, D. J. (2018), "Gaussian variational approximation with sparse precision matrices", *Statistics and Computing*, 28, 259-275.

See Also

extractELBO, fitted.gpr, predict.gpr, plot.gpr, summary.gpr

Examples

```
## Not run:
sdat <- bkmr::SimData()
y <- sdat$y
X <- sdat$X
Z <- sdat$Z

fout <- vbayesGP::gvagpr(y, X, Z, priors = list(lengthscale = 'normal'), covstr = 'diagonal')
plot(fout)
summary(fout)
vbayesGP::extractELBO(fout) # ELBO
## End(Not run)
```

metals

*Urinary metal concentrations and hypertension in NHANES***Description**

Adults aged 20 and over from four cycles of the National Health and Nutrition Examination Survey (2011–2012 through 2017–2018) who were part of the urinary metals subsample and had complete records on the variables below. Concentrations are as reported by the survey and are not transformed; the analysis in the package vignette and in Jo, Ye, Hahn and Lee (2026) takes logarithms and standardizes them.

Usage

metals

Format

A data frame with 5809 rows and 23 columns.

Ba, Cd, Co, Cs, Mo, Mn, Pb, Sb, Sn, Tl, W urinary concentrations of barium, cadmium, cobalt, caesium, molybdenum, manganese, lead, antimony, tin, thallium and tungsten, in ug/L.

sbp, dbp systolic and diastolic blood pressure in mmHg, averaged over the available readings.

meds logical: reported use of antihypertensive medication.

htn hypertension, defined as $sbp \geq 140$, $dbp \geq 90$ or **meds**. Redefine it from **sbp**, **dbp** and **meds** to use another threshold.

age age in years.

sex factor, "male" or "female".

race factor with levels "MexAm", "OthHispanic", "NHWhite", "NHBlack", "NHAsian" and "Other".

pir ratio of family income to the poverty threshold.

bmi body mass index in kg/m^2 .

ucr urinary creatinine in mg/dL.

cotinine serum cotinine in ng/mL.

cycle factor giving the survey cycle.

Source

National Health and Nutrition Examination Survey, cycles 2011–2012, 2013–2014, 2015–2016 and 2017–2018. <https://www.cdc.gov/nchs/nhanes/>

References

Jo S, Ye S, Hahn G, Lee W (2026). Accelerated Bayesian Kernel Machine Regression: A Gaussian Variational Approximation with the Horseshoe Prior. *Statistics and Computing*, **36**, 79.

Examples

```
data(metals)
table(metals$htn)
summary(metals$Pb)
```

`negbinomial`*Negative binomial family for count outcomes*

Description

A family object for the negative binomial distribution with the log link, for use with [gvaggpr](#). The linear predictor is the log mean, as in `stats::glm` and `MASS::glm.nb`, so that `exp(beta)` is a rate ratio.

Usage

```
negbinomial(link = "log")
```

Arguments

`link` the link function. Only "log" is available.

Details

The dispersion is always estimated: overdispersion is the reason to use a negative binomial model in the first place, so there is no option to hold it fixed, and the family therefore takes no dispersion argument. A starting value may be supplied through `inits$theta` in the call to `gvaggpr`; see the examples.

`MASS::negative.binomial(theta = )` is also accepted, with its `theta` taken as a starting value rather than as a fixed dispersion, and a warning to that effect.

Value

an object of class "family".

See Also

[gvaggpr](#)

Examples

```
## Not run:
fit <- gvaggpr(y, X, Z, family = negbinomial())

## with a starting value for the dispersion
fit <- gvaggpr(y, X, Z, family = negbinomial(), inits = list(theta = 5))
## End(Not run)
```

plot.gpmim	<i>Plot Diagnostics for a gpmim Object</i>
------------	--

Description

Provides a plot of the smoothed evidence lower bound (ELBO) against iterations for checking the convergence.

Usage

```
## S3 method for class 'gpmim'
plot(x, ...)
```

Arguments

x	gpmim object, result of gvagpmmim .
...	unused argument.

Author(s)

Seongil Jo

See Also

gvagpmmim

plot.gpr	<i>Plot Diagnostics for a gpr Object</i>
----------	--

Description

Provides a plot of the smoothed evidence lower bound (ELBO) against iterations for checking the convergence.

Usage

```
## S3 method for class 'gpr'
plot(x, from = 1, ...)
```

Arguments

x	gpr object, result of gvagpr or gvaggpr .
from	(positive integer) first iteration to show. The default, 1, plots the whole trace.
...	unused argument.

Details

The bound rises steeply over the first few iterations, often by several orders of magnitude, and on the full range that early climb takes up the whole vertical axis: the plateau where convergence is actually decided is then flat to within a pixel. Use `from` to start the plot after the climb and read the plateau.

Author(s)

Seongil Jo

See Also

gvagpr, gvaggpr

plot.gprfitBivar	<i>Plot bivariate predictor-response function on a new grid of points</i>
------------------	---

Description

Provides a plot of bivariate predictor-response function on a new grid of points

Usage

```
## S3 method for class 'gprfitBivar'
plot(x, ...)
```

Arguments

x	gpr object, result of predictorResponseBivar .
...	unused argument.

Author(s)

Seongil Jo

See Also

gvagpr predictorResponseBivar

plot.gprfitBivarLevels	<i>Plot interactions</i>
------------------------	--------------------------

Description

Provides a plot of the predictor-response function of a single predictor in Z for the second predictor in Z fixed at various quantiles

Usage

```
## S3 method for class 'gprfitBivarLevels'
plot(x, ...)
```

Arguments

x	gpr object, result of predictorResponseBivarLevels .
...	unused argument.

Author(s)

Seongil Jo

See Also

gvagpr predictorResponseBivar

plot.gprfitUnivar	<i>Plot univariate predictor-response function on a new grid of points</i>
-------------------	--

Description

Provides a plot of univariate predictor-response function on a new grid of points

Usage

```
## S3 method for class 'gprfitUnivar'
plot(x, ...)
```

Arguments

x	gpr object, result of predictorResponseUnivar .
...	unused argument.

Author(s)

Seongil Jo

See Also

gvagpr predictorResponseUnivar

plot.risks	<i>Plot Risk Summaries</i>
------------	----------------------------

Description

Provides a plot of risk summaries

Usage

```
## S3 method for class 'risks'
plot(x, ...)
```

Arguments

x	risks object, result of computeOverallRisk, computeSingVarRisk, or computeSingVarInt
...	unused argument

Author(s)

Seongil Jo

See Also

gvagpr computeOverallRisk computeSingVarRisk computeSingVarInt

predict.gpmim

*Extract GPMIM Predicted Values***Description**

predicted is a generic function which extracts predicted values for nonparametric part from an object of class "gpmim"

Usage

```
## S3 method for class 'gpmim'
predict(object, Z_new, nsamples = 1000, ...)
```

Arguments

object	an object of class gpmim.
Z_new	a matrix of new predictor values at which to predict new f, where each row represents a new observation.
nsamples	(positive integer) number of posterior samples. Default value is 1000.
...	additional arguments affecting the predictions produced.

Value

fmean	posterior mean of nonparametric part.
fcov	posterior variance of nonparametric part.

an object of class "gprfit", which has the associated method:
 * plot (i.e., plot.gprfit)

Author(s)

Seongil Jo

See Also

gvagpmmim

predict.gpr

Extract GPR and GGPR Model Predicted Values

Description

predicted is a generic function which extracts predicted values for nonparametric part from an object of class "gpr"

Usage

```
## S3 method for class 'gpr'
predict(object, Z_new, nsamples = 1000, cov = TRUE, ...)
```

Arguments

object	an object of class gpr.
Z_new	a matrix of new predictor values at which to predict new f, where each row represents a new observation.
nsamples	(positive integer) number of posterior samples. Default value is 1000.
cov	logical; if TRUE (the default) the posterior covariance of the new f is returned as well. With many new points this is the dominant cost, so set it to FALSE when only the predictive mean is wanted.
...	additional arguments affecting the predictions produced.

Value

fmean	posterior mean of nonparametric part.
fcov	posterior variance of nonparametric part.

an object of class "gprfit", which has the associated method:
 * plot (i.e., plot.gprfit)

Author(s)

Seongil Jo

See Also

gvagpr, gvaggpr

predictorResponseBivar

Predict the exposure-response function at a new grid of points

Description

Predict the exposure-response function at a new grid of points

Usage

```
predictorResponseBivar(
  fit,
  z.pairs = NULL,
  ngrid = 50,
  q.fixed = 0.5,
  nsamples = 1000,
  min.plot.dist = 0.5,
  center = TRUE,
  verbose = TRUE,
  ...
)
```

Arguments

fit	an object of class gpr
z.pairs	data frame showing which pairs of predictors to plot
ngrid	number of grid points in each dimension
q.fixed	a second quantile at which to compare the estimated f function
nsamples	(positive integer) number of posterior samples. Default value is 1000
min.plot.dist	specifies a minimum distance that a new grid point needs to be from an observed data point in order to compute the prediction; points further than this will not be computed
center	flag for whether to scale the exposure-response function to have mean zero
verbose	TRUE or FALSE: flag of whether to print intermediate output to the screen
...	additional arguments affecting the predictions produced

Value

a long data frame with the name of the first predictor, the name of the second predictor, the value of the first predictor, the value of the second predictor, the posterior mean estimate, and the posterior standard deviation of the estimated exposure response function

References

Bobb J (2023). `_bkmr: Bayesian Kernel Machine Regression_`. R package version 0.2.2, <<https://github.com/jenfb/bkmr>>

Examples

```
## Not run:
## First generate dataset
set.seed(111)
dat <- bkmr::SimData(n = 50, M = 4)
y <- dat$y
Z <- dat$Z
X <- dat$X

set.seed(111)
priors = list(lengthscale = 'horseshoe')
fout <- vbayesGP::gvagpr(y = y, Z = Z, X = X, priors = priors, covstr = 'diagonal')

## Obtain predicted value on new grid of points for each pair of predictors
## Using only a 10-by-10 point grid to make example run quickly
pred.resp.bivar <- vbayesGP::predictorResponseBivar(fit = fout, min.plot.dist = 1, ngrid = 10)
## End(Not run)
```

predictorResponseBivarLevels

Plot cross-sections of the bivariate predictor-response function

Description

Function to plot the f function of a particular variable at different levels (quantiles) of a second variable

Usage

```
predictorResponseBivarLevels(
  object,
  Z = NULL,
  qs = c(0.25, 0.5, 0.75),
  both_pairs = TRUE
)
```

Arguments

object	object obtained from running the function predictorResponseBivar
Z	an n-by-M matrix of predictor variables to be included in nonparametric part.
qs	vector of quantiles at which to fix the second variable
both_pairs	flag indicating whether, if $h(z_1)$ is being plotted for z_2 fixed at different levels, that they should be plotted in the reverse order as well (for $h(z_2)$ at different levels of z_1)

Value

a long data frame with the name of the first predictor, the name of the second predictor, the value of the first predictor, the quantile at which the second predictor is fixed, the posterior mean estimate, and the posterior standard deviation of the estimated exposure response function

References

Bobb J (2023). `_bkmr: Bayesian Kernel Machine Regression_`. R package version 0.2.2, <<https://github.com/jenfb/bkmr>>

Examples

```
## Not run:
## First generate dataset
set.seed(111)
dat <- bkmr::SimData(n = 50, M = 4)
y <- dat$y
Z <- dat$Z
X <- dat$X

set.seed(111)
priors = list(lengthscale = 'horseshoe')
fout <- vbayesGP::gvagpr(y = y, Z = Z, X = X, priors = priors, covstr = 'diagonal')

## Obtain predicted value on new grid of points for each pair of predictors
## Using only a 10-by-10 point grid to make example run quickly
pred.resp.bivar <- vbayesGP::predictorResponseBivar(fit = fout, min.plot.dist = 1, ngrid = 10)
pred.resp.bivar.levels <- vbayesGP::predictorResponseBivarLevels(pred.resp.df = pred.resp.bivar,
  Z = Z, qs = c(0.1, 0.5, 0.9))
## End(Not run)
```

predictorResponseBivarPair

Predict bivariate predictor-response function on a new grid of points

Description

Predict bivariate predictor-response function on a new grid of points

Usage

```
predictorResponseBivarPair(
  fit,
  whichz1 = 1,
  whichz2 = 2,
  whichz3 = NULL,
  prob = 0.5,
  q.fixed = 0.5,
  nsamples = 1000,
  ngrid = 50,
  min.plot.dist = 0.5,
  center = TRUE,
  ...
)
```

Arguments

<code>fit</code>	an object of class <code>gpr</code>
<code>whichz1</code>	vector identifying the first predictor that (column of <code>Z</code>) should be plotted

whichz2	vector identifying the second predictor that (column of Z) should be plotted
whichz3	vector identifying the third predictor that will be set to a pre-specified fixed quantile (determined by prob)
prob	pre-specified quantile to set the third predictor (determined by whichz3); defaults to 0.5 (50th percentile)
q.fixed	a second quantile at which to compare the estimated f function
nsamples	(positive integer) number of posterior samples. Default value is 1000
ngrid	number of grid points in each dimension
min.plot.dist	specifies a minimum distance that a new grid point needs to be from an observed data point in order to compute the prediction; points further than this will not be computed
center	flag for whether to scale the exposure-response function to have mean zero
...	additional arguments affecting the predictions produced

Value

a data frame with value of the first predictor, the value of the second predictor, the posterior mean estimate, and the posterior standard deviation

References

Bobb J (2023). `_bkmr: Bayesian Kernel Machine Regression_`. R package version 0.2.2, <<https://github.com/jenfb/bkmr>>

Examples

```
## Not run:
## First generate dataset
set.seed(111)
dat <- bkmr::SimData(n = 50, M = 4)
y <- dat$y
Z <- dat$Z
X <- dat$X

set.seed(111)
priors = list(lengthscale = 'horseshoe')
fout <- vbayesGP::gvagr(y = y, Z = Z, X = X, priors = priors, covstr = 'diagonal')

## Obtain predicted value on new grid of points
## Using only a 10-by-10 point grid to make example run quickly
pred.resp.bivar12 <- vbayesGP::predictorResponseBivarPair(fit = fout, min.plot.dist = 1, ngrid = 10)
## End(Not run)
```

predictorResponseUnivar

Predict univariate predictor-response function on a new grid of points

Description

Predict univariate predictor-response function on a new grid of points

Usage

```
predictorResponseUnivar(
  fit,
  which.z = 1:ncol(Z),
  ngrid = 50,
  q.fixed = 0.5,
  nsamples = 1000,
  min.plot.dist = Inf,
  center = TRUE,
  z.names = colnames(Z),
  ...
)
```

Arguments

<code>fit</code>	an object of class <code>gpr</code>
<code>which.z</code>	vector identifying which predictors (columns of <code>Z</code>) should be plotted
<code>ngrid</code>	number of grid points to cover the range of each predictor (column in <code>Z</code>)
<code>q.fixed</code>	a second quantile at which to compare the estimated <code>f</code> function
<code>nsamples</code>	(positive integer) number of posterior samples. Default value is 1000
<code>min.plot.dist</code>	specifies a minimum distance that a new grid point needs to be from an observed data point in order to compute the prediction; points further than this will not be computed
<code>center</code>	flag for whether to scale the exposure-response function to have mean zero
<code>z.names</code>	a vector of names of predictors <code>Z</code> . Default values are <code>colnames(Z)</code> .
<code>...</code>	additional arguments affecting the predictions produced.

Value

a long data frame with the predictor name, predictor value, posterior mean estimate, and posterior standard deviation

References

Bobb, J. F., Valeri, L., Claus, H. B., Christiani, D. C., Wright, R. O., Mazumdar, M., Godleski, J. J., and Coull, B. A. (2015). "Bayesian Kernel Machine Regression for Estimating the Health Effects of Multi-Pollutant Mixtures", *Biostatistics*, 16, 493-508.

Bobb J (2023). `_bkmr: Bayesian Kernel Machine Regression_`. R package version 0.2.2, <<https://github.com/jenfb/bkmr>>

Examples

```
## Not run:
## First generate dataset
set.seed(111)
dat <- bkmr::SimData(n = 50, M = 4)
y <- dat$y
Z <- dat$Z
X <- dat$X

set.seed(111)
priors <- list(lengthscale = 'horseshoe')
```

```
fout <- vbayesGP::gvagpr(y = y, Z = Z, X = X, priors = priors, covstr = 'diagonal')
pred.resp.univar <- vbayesGP::predictorResponseUnivar(fout)
## End(Not run)
```

print.gpmim	<i>Print basic summary of gpmim and ggpmim model fit</i>
-------------	--

Description

print method for class "gpmim"

Usage

```
## S3 method for class 'gpmim'
print(x, ...)
```

Arguments

x	an object of class gpmim.
...	unused argument.

Author(s)

Seongil Jo

See Also

gvagpmim

print.gpr	<i>Print basic summary of gpr and ggpr model fit</i>
-----------	--

Description

print method for class "gpr"

Usage

```
## S3 method for class 'gpr'
print(x, ...)
```

Arguments

x	an object of class gpr.
...	unused argument.

Author(s)

Seongil Jo

See Also

gvagpr, gvaggpr

summary.gpmim

*Summarizing gpmim and ggpmim model fits***Description**

summary method for class "gpmim"

Usage

```
## S3 method for class 'gpmim'
summary(
  object,
  q = c(0.025, 0.975),
  digits = 5,
  nsamples = 1000,
  ntuning = 10,
  ...
)
```

Arguments

object	an object of class gpmim
q	quantiles of posterior distribution (credible interval) to show
digits	the number of digits to show when printing
nsamples	(positive integer), number of posterior samples to draw and save, defaults to 1000
ntuning	the number of chosen values of the tuning parameter for variable selection, defaults to 10
...	unused argument.

Author(s)

Seongil Jo

See Also

gvagpmim

summary.gpr

*Summarizing gpr and ggpr model fits***Description**

summary method for class "gpr"

Usage

```
## S3 method for class 'gpr'
summary(
  object,
  q = c(0.025, 0.975),
  digits = 5,
  nsamples = 1000,
  ntuning = 10,
  ...
)
```

Arguments

object	an object of class gpr
q	quantiles of posterior distribution (credible interval) to show
digits	the number of digits to show when printing
nsamples	(positive integer), number of posterior samples to draw and save, defaults to 1000
ntuning	the number of chosen values of the tuning parameter for variable selection, defaults to 10
...	unused argument.

Author(s)

Seongil Jo

See Also

gvagpr, gvaggpr

Index

* datasets

metals, [20](#)

computeOverallRisk, [2](#)

computePPIPs, [3](#)

computeSingVarInt, [4](#)

computeSingVarRisk, [5](#)

extractELB0, [6](#)

extractPostSamps, [7](#)

fitted.gpmim, [7](#)

fitted.gpr, [8](#)

gvaggpr, [9](#), [21](#)

gvagpmim, [13](#)

gvagpr, [16](#)

metals, [20](#)

negbinomial, [21](#)

plot.gpmim, [22](#)

plot.gpr, [22](#)

plot.gprfitBivar, [23](#)

plot.gprfitBivarLevels, [23](#)

plot.gprfitUnivar, [24](#)

plot.risks, [24](#)

predict.gpmim, [25](#)

predict.gpr, [26](#)

predictorResponseBivar, [27](#), [28](#)

predictorResponseBivarLevels, [28](#)

predictorResponseBivarPair, [29](#)

predictorResponseUnivar, [30](#)

print.gpmim, [32](#)

print.gpr, [32](#)

summary.gpmim, [33](#)

summary.gpr, [33](#)